

Tendermint Network Metering Reactor

Charles Dusek

May 1, 2020

Abstract

Decentralized network load balancing requires secure metering. In order to accomplish this a network metering reactor will be designed for the Tendermint consensus engine. This reactor will gather network measurements such as bandwidth, latency and packet loss from a set of neighbors and share those on a BFT CRDT.

1 Introduction

As the number of users increases on an autonomous Byzantine Fault Tolerant decentralized system, load balancing is necessary to meet critical time constraints that dictate user experience. The first step towards load balancing of a decentralized system is generalized network metering. This includes bandwidth, latency, loss, as well as RTT for various signed computational challenges to determine the nodes load.

Decentralized network coordinate methodologies are used to generate estimated maps of networks in a way that minimizes communication, increases fault tolerance and shares the load of computing by co-operative machine learning. The particular method that is used as part of this protocol may be generalized for any network measurement and is node-churn stable by using a self-stabilizing distributed maximum margin matrix factorization algorithm.

Within an adversarial environment, Byzantine node are incentivized to report false network measurements in order receive more rewards by reporting lower lower latency. Byzantine Fault Tolerant consensus allows for verification of tests and agreement on violations of the network metering protocol. Pair-wise network metering relies on nodes interacting and reporting network measurements in a truthful manner. In order to verify the measurements are taken and shared according to the protocol, an added verification step is needed that statistically samples the nodes on the system for faulty protocol implementations.

Each node entrant to the system will be accepted based on their network coordinate error or skew. Any node that is exhibiting faulty behavior will not be accepted into the network and will be reported. Additionally, during each round, nodes will sample from their neighbor set to compare their neighbors coordinates with actual network measurements that are taken during the round. Any node exhibiting out of bounds error between coordinate estimated network measurements and actual measurement will be reported. Finally, reports of accepted nodes exhibiting faulty behavior will be verified through a consensus mechanism that verifies measurements between the the reporting node and the accused node

using third party nodes. If either the faulty node or a byzantine reporter are confirmed by consensus agreement amongst accepted nodes that they are guilty of the accusation the convicted node will be blacklisted.

Algorithm 1: Network Metering Protocol - PLUG To Be Completed

```

Input: Your Input
Output: Your output
Data: Testing set  $x$ 
1  $\sum_{i=1}^{\infty} := 0$  // this is a comment
2 /* Now this is an if...else conditional loop */
3 if Condition 1 then
4   | Do something // this is another comment
5   | if sub-Condition then
6   | | Do a lot
7 else if Condition 2 then
8   | Do Otherwise
9   | /* Now this is a for loop */
10  | for sequence do
11  | | loop instructions
12 else
13 | Do the rest
14 /* Now this is a While loop */
15 while Condition do
16 | Do something

```

Consensus within the network is assured through the Tendermint protocol. Tendermint is a Byzantine Fault Tolerant consensus engine that is used primarily to build blockchain applications. Tendermint uses a gossip protocol to maintain a membership list as well as disseminate messages. Tendermint consensus may tolerate only f faulty or adversarial nodes and proposals are committed by a $2f+1$ vote.

The program structure of Tendermint follows the reactor pattern. Early version of the reactor pattern may be found within the book "Pattern Languages of Program Design" by Jim Coplien and Douglas Schmidt in 1995. An updated version was written by Douglas Schmidt called Reactor: An Object Behavioral Pattern for Demultiplexing and Dispatching Handles for Synchronous Events.

1.1 Command Line Interface

Tendermint uses the Cobra command line interface commander. Cobra provides the following features:

- Easy subcommand-based CLIs: app server, app fetch, etc.
- Fully POSIX-compliant flags (including short and long versions)
- Nested subcommands
- Global, local and cascading flags

- Easy generation of applications and commands with `cobra init appname` and `cobra add cmdname`
- Intelligent suggestions (app server... did you mean app server?)
- Automatic help generation for commands and flags
- Automatic help flag recognition of `-h`, `-help`, etc.
- Automatically generated bash autocomplete for your application
- Automatically generated man pages for your application
- Command aliases so you can change things without breaking them
- The flexibility to define your own help, usage, etc.
- Optional tight integration with viper for 12-factor apps

1.1.1 Main

While you are welcome to provide your own organization, typically a Cobra-based application will follow the following organizational structure:

- appName
 - cmd
 - * add.go
 - * your.go
 - * commands.go
 - * here.go
 - main.go

While most implementations of Cobra use separate main.go and root.go files, Tendermint combines these two files into the main.go and places it in the ./cmd/main.go

```
package main

import (
    "os"
    "path/filepath"

    cmd "github.com/tendermint/tendermint/cmd/tendermint/
        commands"
    "github.com/tendermint/tendermint/cmd/tendermint/commands/
        debug"
    cfg "github.com/tendermint/tendermint/config"
    "github.com/tendermint/tendermint/libs/cli"
    nm "github.com/tendermint/tendermint/node"
)

func main() {
    rootCmd := cmd.RootCmd
    rootCmd.AddCommand(
        cmd.GenValidatorCmd,
        cmd.InitFilesCmd,
        cmd.ProbeUpnpCmd,
        cmd.LiteCmd,
        cmd.ReplayCmd,
        cmd.ReplayConsoleCmd,
        cmd.ResetAllCmd,
        cmd.ResetPrivValidatorCmd,
        cmd.ShowValidatorCmd,
        cmd.TestnetFilesCmd,
        cmd.ShowNodeIDCmd,
        cmd.GenNodeKeyCmd,
        cmd.VersionCmd,
        debug.DebugCmd,
        cli.NewCompletionCmd(rootCmd, true),
    )

    // NOTE:
    // Users wishing to:
    // * Use an external signer for their validators
    // * Supply an in-proc abci app
    // * Supply a genesis doc file from another source
    // * Provide their own DB implementation
    // can copy this file and use something other than the
    // DefaultNewNode function
    nodeFunc := nm.DefaultNewNode

    // Create and start node
    rootCmd.AddCommand(cmd.NewRunNodeCmd(nodeFunc))

    cmd := cli.PrepareBaseCmd(rootCmd, "TM", os.ExpandEnv(
        filepath.Join("$HOME", cfg.DefaultTendermintDir)))
    if err := cmd.Execute(); err != nil {
        panic(err)
    }
}
```

1.1.2 Root

The RootCmd is found within the ./cmd/tendermint/commands/root.go file:

```

// RootCmd is the root command for Tendermint core.
var RootCmd = &cobra.Command{
    Use: "tendermint",
    Short: "Tendermint Core (BFT Consensus) in Go",
    PersistentPreRunE: func(cmd *cobra.Command, args []string)
        (err error) {
            if cmd.Name() == VersionCmd.Name() {
                return nil
            }
            config, err = ParseConfig()
            if err != nil {
                return err
            }
            if config.LogFormat == cfg.LogFormatJSON {
                logger = log.NewTMJSONLogger(log.NewSyncWriter(os.
                    Stdout))
            }
            logger, err = tmflags.ParseLogLevel(config.LogLevel,
                logger, cfg.DefaultLogLevel())
            if err != nil {
                return err
            }
            if viper.GetBool(cli.TraceFlag) {
                logger = log.NewTracingLogger(logger)
            }
            logger = logger.With("module", "main")
            return nil
        },
}

```

1.2 tendermint init

InitFilesCmd initialises a fresh Tendermint Core instance. The InitFileCmd cobra command variable uses the handle "init" and is found within the ./cmd/-tendermint/commands/init.go file:

```
var InitFilesCmd = &cobra.Command{
    Use:   "init",
    Short: "Initialize Tendermint",
    RunE:  initFiles,
}

func initFiles(cmd *cobra.Command, args []string) error {
    return initFilesWithConfig(config)
}

func initFilesWithConfig(config *cfg.Config) error {
    // private validator
    privValKeyFile := config.PrivValidatorKeyFile()
    privValStateFile := config.PrivValidatorStateFile()
    var pv *privval.FilePV
    if tmos.FileExists(privValKeyFile) {
        pv = privval.LoadFilePV(privValKeyFile, privValStateFile)
    }
    logger.Info("Found private validator", "keyFile",
        privValKeyFile,
        "stateFile", privValStateFile)
} else {
    pv = privval.GenFilePV(privValKeyFile, privValStateFile)
    pv.Save()
    logger.Info("Generated private validator", "keyFile",
        privValKeyFile,
        "stateFile", privValStateFile)
}

nodeKeyFile := config.NodeKeyFile()
if tmos.FileExists(nodeKeyFile) {
    logger.Info("Found node key", "path", nodeKeyFile)
} else {
    if _, err := p2p.LoadOrGenNodeKey(nodeKeyFile); err !=
        nil {
        return err
    }
    logger.Info("Generated node key", "path", nodeKeyFile)
}

// genesis file
genFile := config.GenesisFile()
if tmos.FileExists(genFile) {
    logger.Info("Found genesis file", "path", genFile)
} else {
    genDoc := types.GenesisDoc{
        ChainID:      fmt.Sprintf("test-chain-%v", tmrand.
            Str(6)),
        GenesisTime:   tmtime.Now(),
        ConsensusParams: types.DefaultConsensusParams(),
    }
    pubKey, err := pv.GetPubKey()
    if err != nil {
        return errors.Wrap(err, "can't get pubkey")
    }
    genDoc.Validators = []types.GenesisValidator{{
```

```

        Address: pubKey.Address(),
        PubKey:  pubKey,
        Power:   10,
    }}

    if err := genDoc.SaveAs(genFile); err != nil {
        return err
    }
    logger.Info("Generated genesis file", "path", genFile)
}

return nil
}

```

1.3 tendermint node

NewRunNodeCmd returns the command that allows the CLI to start a node. It can be used with a custom PrivValidator and in-process ABCI application. NewRunNodeCmd is found in ./cmd/tendermint/commands/run_node.go:

```

func NewRunNodeCmd(nodeProvider nm.Provider) *cobra.Command
{
    cmd := &cobra.Command{
        Use:   "node",
        Short: "Run the tendermint node",
        RunE: func(cmd *cobra.Command, args []string) error {
            if err := checkGenesisHash(config); err != nil {
                return err
            }

            n, err := nodeProvider(config, logger)
            if err != nil {
                return fmt.Errorf("failed to create node: %w", err)
            }

            if err := n.Start(); err != nil {
                return fmt.Errorf("failed to start node: %w", err)
            }

            logger.Info("Started node", "nodeInfo", n.Switch().
                NodeInfo())

            // Stop upon receiving SIGTERM or CTRL-C.
            tmos.TrapSignal(logger, func() {
                if n.IsRunning() {
                    n.Stop()
                }
            })

            // Run forever.
            select {}
        },
    },
}

AddNodeFlags(cmd)
return cmd
}

```

The `NewRunNodeCmd` function takes the `nodeProvider` as parameter with type `nm.Provider`. The `nm.Provider` type is found in the "node" go package which is the main entry point, where the `Node` struct, which represents a full node, is defined. The `Provider` is imported into the `./cmd/main.go` file as a function whose default is `DefaultNewNode`.

`Provider` takes a `config` and a `logger` and returns a ready to go `Node`. The `nm.Provider` type is found in the `./node/node.go` file:

```
type Provider func(*cfg.Config, log.Logger) (*Node, error)
```

`DefaultNewNode`, which is also found in the `./node/node.go` file, returns a Tendermint node with default settings for the `PrivValidator`, `ClientCreator`, `GenesisDoc`, and `DBProvider`. It implements `NodeProvider`.

```
func DefaultNewNode(config *cfg.Config, logger log.Logger)
(*Node, error) {
    // Generate node PrivKey
    nodeKey, err := p2p.LoadOrGenNodeKey(config.NodeKeyFile())
    if err != nil {
        return nil, err
    }

    // Convert old PrivValidator if it exists.
    oldPrivVal := config.OldPrivValidatorFile()
    newPrivValKey := config.PrivValidatorKeyFile()
    newPrivValState := config.PrivValidatorStateFile()
    if _, err := os.Stat(oldPrivVal); !os.IsNotExist(err) {
        oldPV, err := privval.LoadOldFilePV(oldPrivVal)
        if err != nil {
            return nil, fmt.Errorf("error reading OldPrivValidator
                from %v: %v", oldPrivVal, err)
        }
        logger.Info("Upgrading PrivValidator file",
            "old", oldPrivVal,
            "newKey", newPrivValKey,
            "newState", newPrivValState,
        )
        oldPV.Upgrade(newPrivValKey, newPrivValState)
    }

    return NewNode(config,
        privval.LoadOrGenFilePV(newPrivValKey, newPrivValState),
        nodeKey,
        proxy.DefaultClientCreator(config.ProxyApp, config.ABCI,
            config.DBDir()),
        DefaultGenesisDocProviderFunc(config),
        DefaultDBProvider,
        DefaultMetricsProvider(config.Instrumentation),
        logger,
    )
}
```

The `Node` type defines the highest level interface to a full Tendermint node. It includes all configuration and running services.

```
type Node struct {
    service.BaseService

    // config
```



```

config          *cfg.Config
genesisDoc      *types.GenesisDoc // initial validator set
privValidator types.PrivValidator // local node's
               validator key

// network
transport *p2p.MultiplexTransport
sw         *p2p.Switch // p2p connections
addrBook  pex.AddrBook // known peers
nodeInfo  p2p.NodeInfo
nodeKey   *p2p.NodeKey // our node privkey
isListening bool

// services
eventBus      *types.EventBus // pub/sub for services
stateDB       dbm.DB
blockStore    *store.BlockStore // store the blockchain
               to disk
bcReactor     p2p.Reactor      // for fast-syncing
mempoolReactor *mempl.Reactor  // for gossiping
               transactions
mempool       mempl.Mempool
consensusState *cs.State      // latest consensus state
consensusReactor *cs.Reactor  // for participating in
               the consensus
pexReactor     *pex.Reactor    // for exchanging peer
               addresses
evidencePool   *evidence.Pool // tracking evidence
proxyApp      proxy.AppConns // connection to the
               application
rpcListeners   []net.Listener // rpc servers
txIndexer     txindex.TxIndexer
indexerService *txindex.IndexerService
prometheusSrv *http.Server
}

```

At the heart of Tendermint is the `NewNode` function. The `NewNode` function returns an object of the type `Node` that represents the Tendermint Node. The `NewNode` function is rather lengthy so it will be broken down into sections by function beginning with parameters to end. The `NewNode` function is found within the `./node/node.go` file:

1.3.1 Parameters

```
func NewNode(
    config *cfg.Config,
    privValidator types.PrivValidator,
    nodeKey *p2p.NodeKey,
    clientCreator proxy.ClientCreator,
    genesisDocProvider GenesisDocProvider,
    dbProvider DBProvider,
    metricsProvider MetricsProvider,
    logger log.Logger,
    options ...Option
) (*Node, error) { ...
```

- `config *cfg.Config`
Config is imported into the `./cmd/main.go` file ...
- `privValidator types.PrivValidator`
The `privValidator` is generated in `NewNode` with the function `privval.LoadOrGenFilePV(newPrivValKey, newPrivValState)`
- `nodeKey *p2p.NodeKey`
Pointer to the node's private key
- `clientCreator proxy.ClientCreator`
Proxy is imported into the `./node/node.go` and `clientCreator` is generated with the function `proxy.DefaultClientCreator(config.ProxyApp, config.ABCI, config.DBDir())`
- `genesisDocProvider GenesisDocProvider`
Generated by the `DefaultGenesisDocProviderFunc(config)` from `config.GenesisDocProvider` returns a `GenesisDoc`. It allows the `GenesisDoc` to be pulled from sources other than the filesystem, for instance from a distributed key-value store cluster. `DefaultGenesisDocProviderFunc` returns a `GenesisDocProvider` that loads the `GenesisDoc` from the `config.GenesisFile()` on the filesystem.

```
type GenesisDocProvider func() (*types.GenesisDoc, error)
    )

func DefaultGenesisDocProviderFunc(config *cfg.Config)
    GenesisDocProvider {
    return func() (*types.GenesisDoc, error) {
        return types.GenesisDocFromFile(config.GenesisFile())
    }
    }
```
- `dbProvider DBProvider`
DBProvider takes a `DBContext` and returns instantiated DBs.

```
type DBProvider func(*DBContext) (dbm.DB, error)
```

- metricsProvider MetricsProvider

```
// MetricsProvider returns a consensus, p2p and mempool
Metrics.
type MetricsProvider func(chainID string) (*cs.Metrics,
    *p2p.Metrics, *mempl.Metrics, *sm.Metrics)

// DefaultMetricsProvider returns Metrics build using
Prometheus client library
// if Prometheus is enabled. Otherwise, it returns no-op
Metrics.
func DefaultMetricsProvider(config *cfg.
    InstrumentationConfig) MetricsProvider {
    return func(chainID string) (*cs.Metrics, *p2p.Metrics
        , *mempl.Metrics, *sm.Metrics) {
        if config.Prometheus {
            return cs.PrometheusMetrics(config.Namespace, "
                chain_id", chainID),
                p2p.PrometheusMetrics(config.Namespace, "
                    chain_id", chainID),
                mempl.PrometheusMetrics(config.Namespace, "
                    chain_id", chainID),
                sm.PrometheusMetrics(config.Namespace, "chain_id
                    ", chainID)
        }
        return cs.NopMetrics(), p2p.NopMetrics(), mempl.
            NopMetrics(), sm.NopMetrics()
    }
}
```

- logger log.Logger
- options ...Option

1.3.2 Initialize Databases

The blockStore and stateDB are generated in the following function:

```
blockStore, stateDB, err := initDBs(config, dbProvider)
if err != nil {
    return nil, err
}
```

The initDBs(config, dbProvider) function is also found in the ./node/node.go file:

```
func initDBs(config *cfg.Config, dbProvider DBProvider) (
    blockStore *store.BlockStore, stateDB dbm.DB, err error)
{
    var blockStoreDB dbm.DB
    blockStoreDB, err = dbProvider(&DBContext{"blockstore",
        config})
    if err != nil {
        return
    }
    blockStore = store.NewBlockStore(blockStoreDB)

    stateDB, err = dbProvider(&DBContext{"state", config})
}
```

```

    if err != nil {
        return
    }

    return
}

```

The DBContext struct type, dbProvider func type, and the DefaultDBProvider function are also declared in the ./node/node.go file:

```

type DBContext struct {
    ID      string
    Config  *cfg.Config
}

type DBProvider func(*DBContext) (dbm.DB, error)

func DefaultDBProvider(ctx *DBContext) (dbm.DB, error) {
    dbType := dbm.BackendType(ctx.Config.DBBackend)
    return dbm.NewDB(ctx.ID, dbType, ctx.Config.DBDir()), nil
}

```

Todo: Fill out the DB portion here starting with dbm.BackendTye()

1.3.3 Initialize State

The state and genDoc are generated with the following function:

```

state, genDoc, err := LoadStateFromDBOrGenesisDocProvider(
    stateDB, genesisDocProvider)
if err != nil {
    return nil, err
}

```

LoadStateFromDBOrGenesisDocProvider attempts to load the state from the database, or creates one using the given genesisDocProvider and persists the result to the database. On success this also returns the genesis doc loaded through the given provider. It is also found within the ./node/node.go file

```

func LoadStateFromDBOrGenesisDocProvider(
    stateDB dbm.DB,
    genesisDocProvider GenesisDocProvider,
) (sm.State, *types.GenesisDoc, error) {
    // Get genesis doc
    genDoc, err := loadGenesisDoc(stateDB)
    if err != nil {
        genDoc, err = genesisDocProvider()
        if err != nil {
            return sm.State{}, nil, err
        }
        // save genesis doc to prevent a certain class of user
        errors (e.g. when it
        // was changed, accidentally or not). Also good for
        audit trail.
        saveGenesisDoc(stateDB, genDoc)
    }
    state, err := sm.LoadStateFromDBOrGenesisDoc(stateDB,
        genDoc)
    if err != nil {
        return sm.State{}, nil, err
    }
    return state, genDoc, nil
}

```

If state already exists with the stateDB then the following function will load it:

```
// panics if failed to unmarshal bytes
func loadGenesisDoc(db dbm.DB) (*types.GenesisDoc, error) {
    b, err := db.Get(genesisDocKey)
    if err != nil {
        panic(err)
    }
    if len(b) == 0 {
        return nil, errors.New("genesis doc not found")
    }
    var genDoc *types.GenesisDoc
    err = cdc.UnmarshalJSON(b, &genDoc)
    if err != nil {
        panic(fmt.Sprintf("Failed to load genesis doc due to
            unmarshaling error: %v (bytes: %X)", err, b))
    }
    return genDoc, nil
}
```

This function saves the genesis doc to prevent a certain class of user errors (e.g. when it was changed, accidentally or not). Also good for audit trail.

```
// panics if failed to marshal the given genesis document
func saveGenesisDoc(db dbm.DB, genDoc *types.GenesisDoc) {
    b, err := cdc.MarshalJSON(genDoc)
    if err != nil {
        panic(fmt.Sprintf("Failed to save genesis doc due to
            marshaling error: %v", err))
    }
    db.SetSync(genesisDocKey, b)
}
```

1.3.4 Create Proxy and ABCI Connections

Create the proxyApp and establish connections to the ABCI app (consensus, mempool, query).

```
proxyApp, err := createAndStartProxyAppConns(clientCreator,
    logger)
if err != nil {
    return nil, err
}
```

The ClientCreator type is found within the ./proxy/client.go file:

```
type ClientCreator interface {
    NewABCIClient() (abci.Client, error)
}
```

The DefaultNewNode function inputs the proxy.DefaultClientCreator as the clientCreator parameter. The DefaultClientCreator is found within the ./proxy/client.go file:

```
func DefaultClientCreator(addr, transport, dbDir string)
    ClientCreator {
    switch addr {
    case "counter":
        return NewLocalClientCreator(counter.NewApplication(
            false))
    case "counter_serial":
```

```

        return NewLocalClientCreator(counter.NewApplication(true
        ))
    case "kvstore":
        return NewLocalClientCreator(kvstore.NewApplication())
    case "persistent_kvstore":
        return NewLocalClientCreator(kvstore.
            NewPersistentKVStoreApplication(dbDir))
    case "noop":
        return NewLocalClientCreator(types.NewBaseApplication())
    default:
        mustConnect := false // loop retrying
        return NewRemoteClientCreator(addr, transport,
            mustConnect)
    }
}

```

Add parameters in Config file here

The helper function for Proxy creation is below:

```

func createAndStartProxyAppConns(clientCreator proxy.
    ClientCreator, logger log.Logger) (proxy.AppConns, error)
{
    proxyApp := proxy.NewAppConns(clientCreator)
    proxyApp.SetLogger(logger.With("module", "proxy"))
    if err := proxyApp.Start(); err != nil {
        return nil, fmt.Errorf("error starting proxy app
            connections: %v", err)
    }
    return proxyApp, nil
}

```

NewABCIClient returns newly connected client that is also found within the ./proxy/client.go file:

```

func (r *remoteClientCreator) NewABCIClient() (abcicli.
    Client, error) {
    remoteApp, err := abcicli.NewClient(r.addr, r.transport, r
        .mustConnect)
    if err != nil {
        return nil, errors.Wrap(err, "Failed to connect to proxy
            ")
    }
    return remoteApp, nil
}

```

NewClient returns a new ABCI client of the specified transport type. It returns an error if the transport is not "socket" or "grpc". The NewClient function is found in the ./abci/client.go file.

```

func NewClient(addr, transport string, mustConnect bool) (
    client Client, err error) {
    switch transport {
    case "socket":
        client = NewSocketClient(addr, mustConnect)
    case "grpc":
        client = NewGRPCClient(addr, mustConnect)
    default:
        err = fmt.Errorf("unknown abci transport %s", transport)
    }
    return
}

```

1.3.5 Event Bus and Indexer Service

EventBus and IndexerService must be started before the handshake because we might need to index the txs of the replayed block as this might not have happened when the node stopped last time (i.e. the node stopped after it saved the block but before it indexed the txs, or, endblocker panicked)

```
eventBus, err := createAndStartEventBus(logger)
if err != nil {
    return nil, err
}

// Transaction indexing
indexerService, txIndexer, err :=
    createAndStartIndexerService(config, dbProvider, eventBus,
    , logger)
if err != nil {
    return nil, err
}
```

The createAndStartEventBus(logger) function is found within the ./node/node.go file:

```
func createAndStartEventBus(logger log.Logger) (*types.
    EventBus, error) {
    eventBus := types.NewEventBus()
    eventBus.SetLogger(logger.With("module", "events"))
    if err := eventBus.Start(); err != nil {
        return nil, err
    }
    return eventBus, nil
}
```

EventBus is a common bus for all events going through the system. All calls are proxied to underlying pubsub server. All events must be published using EventBus to ensure correct data types. The EventBus functionality below is found within the ./types/event_bus.go file:

```
type EventBus struct {
    service.BaseService
    pubsub *tmpubsub.Server
}

// NewEventBus returns a new event bus.
func NewEventBus() *EventBus {
    return NewEventBusWithBufferCapacity(defaultCapacity)
}

// NewEventBusWithBufferCapacity returns a new event bus
// with the given buffer capacity.
func NewEventBusWithBufferCapacity(cap int) *EventBus {
    // capacity could be exposed later if needed
    pubsub := tmpubsub.NewServer(tmpubsub.BufferCapacity(cap))
    b := &EventBus{pubsub: pubsub}
    b.BaseService = *service.NewBaseService(nil, "EventBus", b)
    return b
}

func (b *EventBus) SetLogger(l log.Logger) {
    b.BaseService.SetLogger(l)
    b.pubsub.SetLogger(l.With("module", "pubsub"))
}
```

Key value indexer is default. It is the simplest possible indexer backed by key-value storage which defaults to LevelDB.

```
// TxIndexConfig
// Remember that Event has the following structure:
// type: [
//   key: value,
//   ...
// ]
//
// CompositeKeys are constructed by 'type.key'
// TxIndexConfig defines the configuration for the
// transaction indexer,
// including composite keys to index.
type TxIndexConfig struct {
    // What indexer to use for transactions
    //
    // Options:
    //   1) "null"
    //   2) "kv" (default) - the simplest possible indexer,
    //      backed by key-value storage (defaults to levelDB;
    //      see DBBackend).
    Indexer string 'mapstructure:"indexer"'

    // Comma-separated list of compositeKeys to index (by
    // default the only key is "tx.hash")
    //
    // You can also index transactions by height by adding "tx
    // .height" key here.
    //
    // It's recommended to index only a subset of keys due to
    // possible memory
    // bloat. This is, of course, depends on the indexer's DB
    // and the volume of
    // transactions.
    IndexKeys string 'mapstructure:"index_keys"'

    // When set to true, tells indexer to index all
    // compositeKeys (predefined keys:
    // "tx.hash", "tx.height" and all keys from DeliverTx
    // responses).
    //
    // Note this may be not desirable (see the comment above).
    // IndexKeys has a
    // precedence over IndexAllKeys (i.e. when given both,
    // IndexKeys will be
    // indexed).
    IndexAllKeys bool 'mapstructure:"index_all_keys"'
}

// DefaultTxIndexConfig returns a default configuration for
// the transaction indexer.
func DefaultTxIndexConfig() *TxIndexConfig {
    return &TxIndexConfig{
        Indexer:      "kv",
        IndexKeys:    "",
        IndexAllKeys: false,
    }
}

// TestTxIndexConfig returns a default configuration for the
// transaction indexer.
```



```
func TestTxIndexConfig() *TxIndexConfig {  
    return DefaultTxIndexConfig()  
}
```

1.3.6 Handshaker

Handshaker calls RequestInfo, sets the AppVersion on the state, and replays any blocks as necessary to sync tendermint with the app.

```
consensusLogger := logger.With("module", "consensus")  
if err := doHandshake(stateDB, state, blockStore, genDoc,  
    eventBus, proxyApp, consensusLogger); err != nil {  
    return nil, err  
}
```

1.4 Communication

There are three forms of communication (e.g., requests, responses, connections) that can happen in Tendermint Core

- Internode: communication between a node and other peers. This kind of communication happens over TCP or HTTP.
- Intranode: communication within the node itself (i.e., between reactors or other components). These are typically function or method calls, or occasionally happen through an event bus.
- Client: communication between a client (like a wallet or a browser) and a node on the network.

2 Tendermint Switch Structure

The Tendermint switch handles the incoming traffic from the P2P network. It is the interface between the node and the rest of the network and routes messages.

- Responsible for routing connections between peers
- Notably: only handles TCP connections; RPC/HTTP is separate
- Is a dependency for every reactor; all reactors expose a function ‘setSwitch’
- Holds onto channels (channels on the TCP connection–NOT Go channels) and uses them to route
- Is a global object, with a global namespace for messages
- Similar functionality to libp2p

2.1 NewSwitch

The NewSwitch function can be found in p2p/switch.go with the code as follows:

```
// NewSwitch creates a new Switch with the given config.
func NewSwitch(
    cfg *config.P2PConfig,
    transport Transport,
    options ...SwitchOption,
) *Switch {
    sw := &Switch{
        config:          cfg,
        reactors:         make(map[string]Reactor),
        chDescs:         make([]*conn.ChannelDescriptor, 0),
        reactorsByCh:     make(map[byte]Reactor),
        peers:           NewPeerSet(),
        dialing:         cmap.NewCMap(),
        reconnecting:    cmap.NewCMap(),
        metrics:         NopMetrics(),
        transport:       transport,
        filterTimeout:   defaultFilterTimeout,
        persistentPeersAddrs: make([]*NetAddress, 0),
        unconditionalPeerIDs: make(map[ID]struct{}),
    }

    // Ensure we have a completely undeterministic PRNG.
    sw.rng = rand.NewRand()

    sw.BaseService = *service.NewBaseService(nil, "P2P Switch", sw)

    for _, option := range options {
        option(sw)
    }

    return sw
}
```

2.2 createSwitch

Core Tendermint reactors are added to the switch in node/node.go

```
func createSwitch(config *cfg.Config,
    transport p2p.Transport,
    p2pMetrics *p2p.Metrics,
    peerFilters []p2p.PeerFilterFunc,
    mempoolReactor *mempl.Reactor,
    bcReactor p2p.Reactor,
    consensusReactor *consensus.Reactor,
    evidenceReactor *evidence.Reactor,
    nodeInfo p2p.NodeInfo,
    nodeKey *p2p.NodeKey,
    p2pLogger log.Logger) *p2p.Switch {

    sw := p2p.NewSwitch(
        config.P2P,
        transport,
        p2p.WithMetrics(p2pMetrics),
        p2p.SwitchPeerFilters(peerFilters...),
    )
    sw.SetLogger(p2pLogger)
    sw.AddReactor("MEMPOOL", mempoolReactor)
    sw.AddReactor("BLOCKCHAIN", bcReactor)
    sw.AddReactor("CONSENSUS", consensusReactor)
    sw.AddReactor("EVIDENCE", evidenceReactor)

    sw.SetNodeInfo(nodeInfo)
    sw.SetNodeKey(nodeKey)

    p2pLogger.Info("P2P Node ID", "ID", nodeKey.ID(), "file",
        config.NodeKeyFile())
    return sw
}
```

2.3 addReactors

2.4 CustomReactors

CustomReactors allows you to add custom reactors (name -> p2p.Reactor) to the node's Switch.

WARNING: using any name from the below list of the existing reactors will result in replacing it with the custom one. - MEMPOOL - BLOCKCHAIN - CONSENSUS - EVIDENCE - PEX

```
func CustomReactors(reactors map[string]p2p.Reactor) Option
{
    return func(n *Node) {
        for name, reactor := range reactors {
            if existingReactor := n.sw.Reactor(name);
                existingReactor != nil {
                n.sw.Logger.Info("Replacing existing reactor with a
                    custom one",
                    "name", name, "existing", existingReactor, "custom
                        ", reactor)
                n.sw.RemoveReactor(name, existingReactor)
            }
            n.sw.AddReactor(name, reactor)
        }
    }
}
```

The addReactor function may be found in the p2p directory inside the switch.go file as follows:

```
// AddReactor adds the given reactor to the switch.
// NOTE: Not goroutine safe.
func (sw *Switch) AddReactor(name string, reactor Reactor)
    Reactor {
    for _, chDesc := range reactor.GetChannels() {
        chID := chDesc.ID
        // No two reactors can share the same channel.
        if sw.reactorsByCh[chID] != nil {
            panic(
                fmt.Sprintf(
                    "Channel %X has multiple reactors %v & %v",
                    chID,
                    sw.reactorsByCh[chID],
                    reactor
                )
            )
        }
        sw.chDescs = append(sw.chDescs, chDesc)
        sw.reactorsByCh[chID] = reactor
    }
    sw.reactors[name] = reactor
    reactor.SetSwitch(sw)
    return reactor
}
```

3 Tendermint Reactor Structure

The base Reactor type may be found in the p2p directory inside the base_reactor.go file.

”Reactor is responsible for handling incoming messages on one or more channels. Switch calls GetChannels when reactor is added to it. When a new peer joins our node, InitPeer and AddPeer are called. RemovePeer is called when the peer is stopped. Receive is called when a message is received on a channel associated with this reactor. Peer Send or Peer TrySend should be used to send the message to a peer.”

```
type Reactor interface {
    service.Service // Start, Stop

    SetSwitch(*Switch)

    GetChannels() []*conn.ChannelDescriptor

    InitPeer(peer Peer) Peer

    AddPeer(peer Peer)

    RemovePeer(peer Peer, reason interface{})

    Receive(chID byte, peer Peer, msgBytes []byte)
}

//-----
```

- `SetSwitch(*Switch)`

`SetSwitch` allows setting a switch. Every reactor holds a pointer to the global switch and the switch holds a pointer to every reactor.

- `GetChannels() []*conn.ChannelDescriptor`

`GetChannels` returns the list of `MConnection.ChannelDescriptor`. Make sure that each ID is unique across all the reactors added to the switch.

- `InitPeer(peer Peer) Peer`

`InitPeer` is called by the switch before the peer is started. Use it to initialize data for the peer (e.g. peer state).

NOTE: The switch won't call `AddPeer` nor `RemovePeer` if it fails to start the peer. Do not store any data associated with the peer in the reactor itself unless you don't want to have a state, which is never cleaned up.

- `AddPeer(peer Peer)`

`AddPeer` is called by the switch after the peer is added and successfully started. Use it to start goroutines communicating with the peer.

- `RemovePeer(peer Peer, reason interface)`

`RemovePeer` is called by the switch when the peer is stopped (due to error or other reason).

- `Receive(chID byte, peer Peer, msgBytes []byte)`

`Receive` is called by the switch when `msgBytes` is received from the peer.

NOTE: reactor can not keep `msgBytes` around after `Receive` completes without copying.

CONTRACT: `msgBytes` are not nil.

3.1 Mock Reactor

The mock reactor file is located in the p2p/mock/reactor.go file

```
package mock

import (
    "github.com/tendermint/tendermint/libs/log"
    "github.com/tendermint/tendermint/p2p"
    "github.com/tendermint/tendermint/p2p/conn"
)

type Reactor struct {
    p2p.BaseReactor
}

func NewReactor() *Reactor {
    r := &Reactor{}
    r.BaseReactor = *p2p.NewBaseReactor("Mock-PEX", r)
    r.SetLogger(log.TestingLogger())
    return r
}

func (r *Reactor) GetChannels() []*conn.ChannelDescriptor
    { return []*conn.ChannelDescriptor{} }
func (r *Reactor) AddPeer(peer p2p.Peer)
    {}
func (r *Reactor) RemovePeer(peer p2p.Peer, reason interface
    {}){}
func (r *Reactor) Receive(chID byte, peer p2p.Peer, msgBytes
    []byte) {}
```

3.2 Existing Tendermint Reactors

Tendermint currently contains five reactors in the